

Gradient Descent Process

Gradient descent is an algorithm that allows you to fit any kind of model to a set of points with any kind of measure of error. This document will walk you through the basic steps of how to do it. As a motivating example, we will refer to the function $y = ax^2 + bx + c$, but know that you will need to change the computations for different types of functions you consider.

Outline

1. **Graph the data and assess potential types of functions** that could fit your dataset.
 - a. Here, we will assume that your visual inspection is a U shape, so you choose a degree 2 polynomial.
2. **Split the data into training, validation, and testing sets**
 - a. Typically, your validation and testing sets are as much as 20% of your data and as little as 5%.
 - b. In the example, we first remove the testing set, then the validation set. This is because you will use both the training and validation to construct your final model

```
X_train_val, X_test, y_train_val, y_test = train_test_split(X, y, test_size=0.1, random_state=0)
X_train, X_val, y_train, y_val = train_test_split(X_train_val, y_train_val, test_size=0.1, random_state=0)
```

3. **Select your error function**
 - a. Mean squared error (MSE) and mean absolute error (MAE) are typical. Mean absolute error is favored in this class.
 - b. The function for MAE is

$$E = \frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i|$$

4. **Compute the partial derivatives** of your error function with respect to the parameters of your model:
 - a. In our model, we will look at the partial derivatives of E with respect to a, b, c .
 - b. For example, for MAE with our model, the derivative with respect to a is:

$$E_a = \frac{1}{n} \sum_{i=1}^n \text{sign}(y_i - (ax^2 + bx + c)) \cdot x^2$$

5. **Pick initial values for your parameters** (coefficient)
 - a. This means you pick an initial value for a, b, c

- b. A great way to do this is plot your guesses on top of your data play around with values until you get something close-ish to your data.
- 6. **Pick initial values for your hyperparameters.**
 - a. **Epochs** – these are the number of iterations (times in the loop you will do). 1000 is usually more than enough.
 - b. **Learning rate** – this is what you will multiply your derivatives by. If you are unsure, a small learning rate, like 0.001, is a good place to start.
- 7. **Perform gradient descent on your training set; evaluate with your validation set.**
 - a. Compute predictions. (\hat{y}) with the current value of the parameters
 - b. Compute the derivatives with the current value of the parameters
 - c. Update the parameters using the gradient descent formula
 - d. Compute the loss/error of the training and the validation sets
- 8. **Repeat this while tweaking your hyperparameters until you see convergence.**
 - a. You want to see your errors get small and level out.
 - b. You also don't want to see your training errors to be much lower than your validation errors.
 - c. Things you can tweak in this process:
 - i. The learning rate
 - ii. Initial parameters
 - d. You may decide to give up on the model and try a new one. If your errors are just consistently bad, it might be best to give up.
 - i. Other things to try:
 - 1. Scale your data to be between -1 and 1.
 - 2. Try a different measure of error.
- 9. **Plot the loss for the validation and training set over the number of iteration/epochs.**
 - a. This will help you assess how many epochs you should have/
- 10. Based on your previous work, **select your hyperparameters.**
 - a. Number of epochs shouldn't be as high as what you did previously. You want a number that happens right when the errors level off but you don't want training to be significantly. Better than validation.
 - b. The learning rate should be small enough to ensure a smooth convergence.
 - c. Your initial parameters should allow you to be able to find good parameters in the end.
- 11. **Perform gradient descent with your selected hyperparameters using your training and validation set combined.**
- 12. Plot the actual data points and the model's prediction.
- 13. **Evaluate the loss using your testing function.** This will tell you how well your model performs.